

# CS 178 HW 2

January 27, 2017

## 1 Kevin Phung, 47881547

### 1.1 Problem 1: Linear Regression

```
In [1]: import numpy as np
import matplotlib.pyplot as plt
import mltools as ml

# Problem 1: Linear Regression


# 1a
data = np.genfromtxt("data/curve80.txt", delimiter=None)

X = data[:,0]
X = X[:,np.newaxis]
Y = data[:,1]
Xtr,Xte,Ytr,Yte = ml.splitData(X,Y,0.75)

# 1b
lr = ml.linear.linearRegress(Xtr,Ytr)
xs = np.linspace(0,10,200)
xs = xs[:,np.newaxis]
ys = lr.predict(xs)

plt.scatter(Xtr,Ytr)
plt.plot(xs,ys,"g")
ax = plt.axis()
plt.show()
print("Linear Regression Coefficients: "+str(lr.theta))
mseTe = lr.mse(Xte,Yte)
mseTr = lr.mse(Xtr,Ytr)

print("Mean Square Error of Test Data: "+str(mseTe))
print("Mean Square Error of Training Data: "+str(mseTr))
```



Linear Regression Coefficients:  $\begin{bmatrix} -2.82765049 & 0.83606916 \end{bmatrix}$   
Mean Square Error of Test Data: 2.24234920301  
Mean Square Error of Training Data: 1.12771195561

```
In [2]: # 1c
degrees = [1,3,5,7,10,18]
mseTeList = []
mseTrList = []

for degree in degrees:

    XtrP = ml.transforms.fpoly(Xtr,degree,bias=False)

    XtrP,params = ml.transforms.rescale(XtrP)

    lr = ml.linear.linearRegress(XtrP,Ytr)

    XteP,params = ml.transforms.rescale(ml.transforms.fpoly(Xte,degree,
                                                                False),params)

    Phi = lambda X: ml.transforms.rescale(ml.transforms.fpoly(X,degree,
                                                                False),params)

    YhatTrain = lr.predict(Phi(xs))
    YhatTest = lr.predict(Phi(Xte))
```

```


plt.plot(xs, YhatTrain, "g")
plt.scatter(Xtr, Ytr)
plt.axis(ax)
plt.title("Degree = "+str(degree))
plt.show()



mseTe = lr.mse(Phi(Xte), Yte)
mseTr = lr.mse(Phi(Xtr), Ytr)



mseTeList.append(mseTe)
mseTrList.append(mseTr)

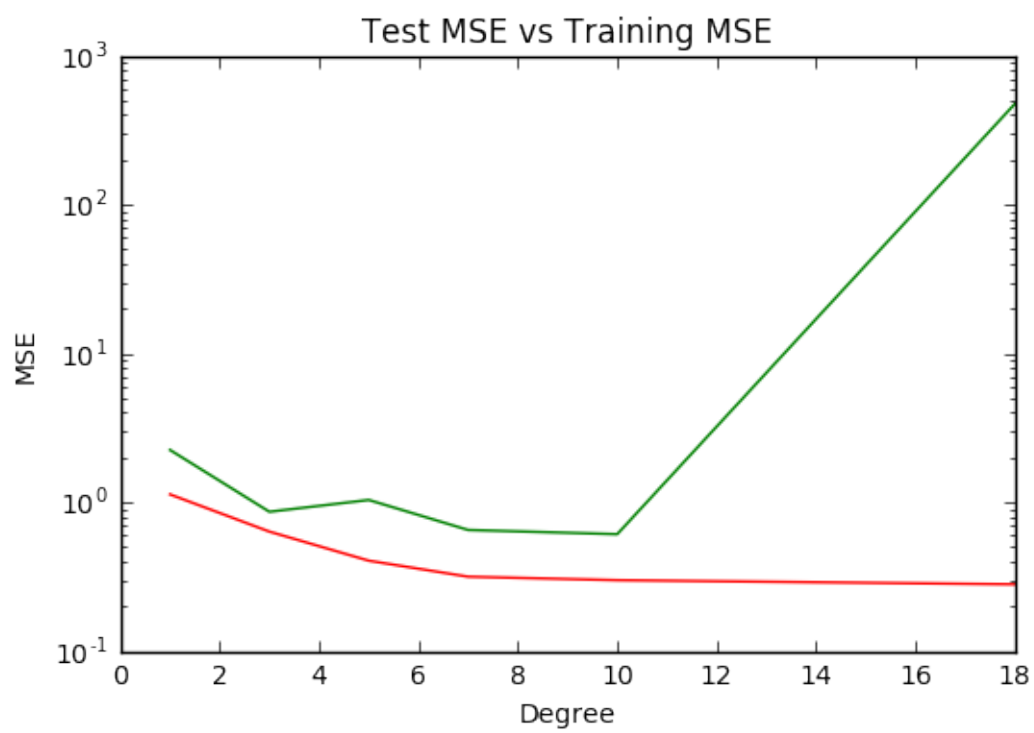

plt.title("Test MSE vs Training MSE")
plt.ylabel("MSE")
plt.xlabel("Degree")
plt.semilogy(degrees, mseTeList, 'g-', degrees, mseTrList, 'r-')
plt.show()
print("Green = Test Data, Red = Training Data")

```









Green = Test Data, Red = Training Data

## 1.2 Problem 2: Cross-Validation


In [3]: *#Problem 2: Cross-validation*

```
nFolds = 5;
J = []

for degree in degrees:
    XtrP = ml.transforms.fpoly(Xtr,degree,bias=False)
    temp = [0,0,0,0,0]
    for iFold in range(nFolds):
        Xti,Xvi,Yti,Yvi = ml.crossValidate(XtrP,Ytr,nFolds,iFold)
        learner = ml.linear.linearRegress(Xti,Yti)
        temp[iFold] = (learner.mse(Xvi,Yvi))
    J.append(np.mean(temp))

plt.title("Cross-Validation vs Actual Test Data")
plt.ylabel("MSE")
plt.xlabel("Degree")
plt.semilogy(degrees,J,'g-',degrees,mseTeList,'r-')
plt.show()
print("Green = Cross-Validation, Red = Actual Test Data")

'''
According to the graph, degree 5 has the minimum cross-validation error.
The MSE estimated from cross-validation is lower compared to the actual
test data on higher degrees. Starting from around 10 degrees the MSE of
the actual test data is substantially higher than the cross-validation.
'''
```



Green = Cross-Validation, Red = Actual Test Data