

CS 178 HW 2

January 27, 2017

1 Kevin Phung, 47881547

1.1 Problem 1: Linear Regression

```
In [1]: import numpy as np
import matplotlib.pyplot as plt
import mltools as ml

# Problem 1: Linear Regression

# 1a
data = np.genfromtxt("data/curve80.txt", delimiter=None)


X = data[:,0]
X = X[:,np.newaxis]
Y = data[:,1]
Xtr,Xte,Ytr,Yte = ml.splitData(X,Y,0.75)

# 1b
lr = ml.linear.linearRegress(Xtr,Ytr)
xs = np.linspace(0,10,200)
xs = xs[:,np.newaxis]
ys = lr.predict(xs)

print("Green = Test Data, Red = Training Data")
plt.scatter(Xtr,Ytr,color = "red")
plt.scatter(Xte,Yte,color = "green")
plt.plot(xs,ys,)
ax = plt.axis()
plt.show()
print("Linear Regression Coefficients: "+str(lr.theta))
mseTe = lr.mse(Xte,Yte)
mseTr = lr.mse(Xtr,Ytr)

print("Mean Square Error of Test Data: "+str(mseTe))
print("Mean Square Error of Training Data: "+str(mseTr))
```

Green = Test Data, Red = Training Data



```
Linear Regression Coefficients: [[-2.82765049  0.83606916]]
Mean Square Error of Test Data: 2.24234920301
Mean Square Error of Training Data: 1.12771195561
```

```
In [2]: # 1c
degrees = [1,3,5,7,10,18]
mseTeList = []
mseTrList = []

for degree in degrees:

    XtrP = ml.transforms.fpoly(Xtr,degree,bias=False)

    XtrP,params = ml.transforms.rescale(XtrP)

    lr = ml.linear.linearRegress(XtrP,Ytr)

    XteP,params = ml.transforms.rescale(ml.transforms.fpoly(Xte,degree,
                                                                False),params)

    Phi = lambda X: ml.transforms.rescale(ml.transforms.fpoly(X,degree,
                                                                False),params)[0]
```

```

YhatTrain = lr.predict(Phi(xs))
YhatTest = lr.predict(Phi(Xte))


plt.plot(xs, YhatTrain)
plt.scatter(Xtr, Ytr, color = "red")
plt.scatter(Xte, Yte, color = "green")
plt.axis(ax)
plt.title("Degree = "+str(degree))
plt.show()



mseTe = lr.mse(Phi(Xte), Yte)
mseTr = lr.mse(Phi(Xtr), Ytr)



mseTeList.append(mseTe)
mseTrList.append(mseTr)



plt.title("Test MSE vs Training MSE")
plt.ylabel("MSE")
plt.xlabel("Degree")
plt.semilogy(degrees, mseTeList, 'g-', degrees, mseTrList, 'r-')
plt.show()
print("Green = Test Data, Red = Training Data")

```









Green = Test Data, Red = Training Data

1.2 Problem 2: Cross-Validation

In [3]: *#Problem 2: Cross-validation*

```
nFolds = 5;
J = []

for degree in degrees:

    temp = [0,0,0,0,0]
    for iFold in range(nFolds):
        Xti,Xvi,Yti,Yvi = ml.crossValidate(Xtr,Ytr,nFolds,iFold)

        XtiP = ml.transforms.fpoly(Xti,degree,bias=False)


        XtiP,params = ml.transforms.rescale(XtiP)

        XtiP,params = ml.transforms.rescale(ml.transforms.fpoly(Xti,degree,
                                                                    False),params)
        Phi = lambda X: ml.transforms.rescale(ml.transforms.fpoly(X,degree,
                                                                    False),params)[0]

        learner = ml.linear.linearRegress(XtiP,Yti)
        temp[iFold] = (learner.mse(Phi(Xvi),Yvi))
    J.append(np.mean(temp))

plt.title("Cross-Validation vs Actual Test Data")
plt.ylabel("MSE")
plt.xlabel("Degree")
plt.semilogy(degrees,J,'b-',degrees,mseTeList,'g-')
plt.show()
print("Blue = Cross-Validation, Green = Actual Test Data")

'''
According to the graph, degree 5 has the minimum cross-validation error.
The MSE estimated from cross-validation is higher compared to the actual
test data on higher degrees. Starting from around 10 degrees and
higher the Cross-Validation data begins to have a much higher MSE
and is overfitting earlier. This is why at lower degrees
Cross-Validation has lower MSE but it quickly raises to high levels
faster than the Test Data.
'''
```



Blue = Cross-Validation, Green = Actual Test Data